

Enhancing Cyber Security with Graph Semantics and KQL

YellowHat 2026

-Henning Rauch-

YELLOWHAT

HARD HATS FOR CYBER THREATS

Elite

VirtualMetric

Gold



glueck■kanja



bnevinn

baseVISION

water
IT Security & Defense



Rubicon

Bronze



SITS

TRAXION

MOSADEX E-HEALTH

valid

INSPARK

yellow arrow

DELTA-N

ENDOR LABS

yellowhat.live



About Me

Current status: Under the radar

Background:

- Former Principal Product Manager at Microsoft
- Part of Kusto Detective Agency (Prof. Smoke)
- LinkedIn: <https://www.linkedin.com/in/henning-rauch-adx/>

Graphs at Microsoft with a strong bias towards Kusto

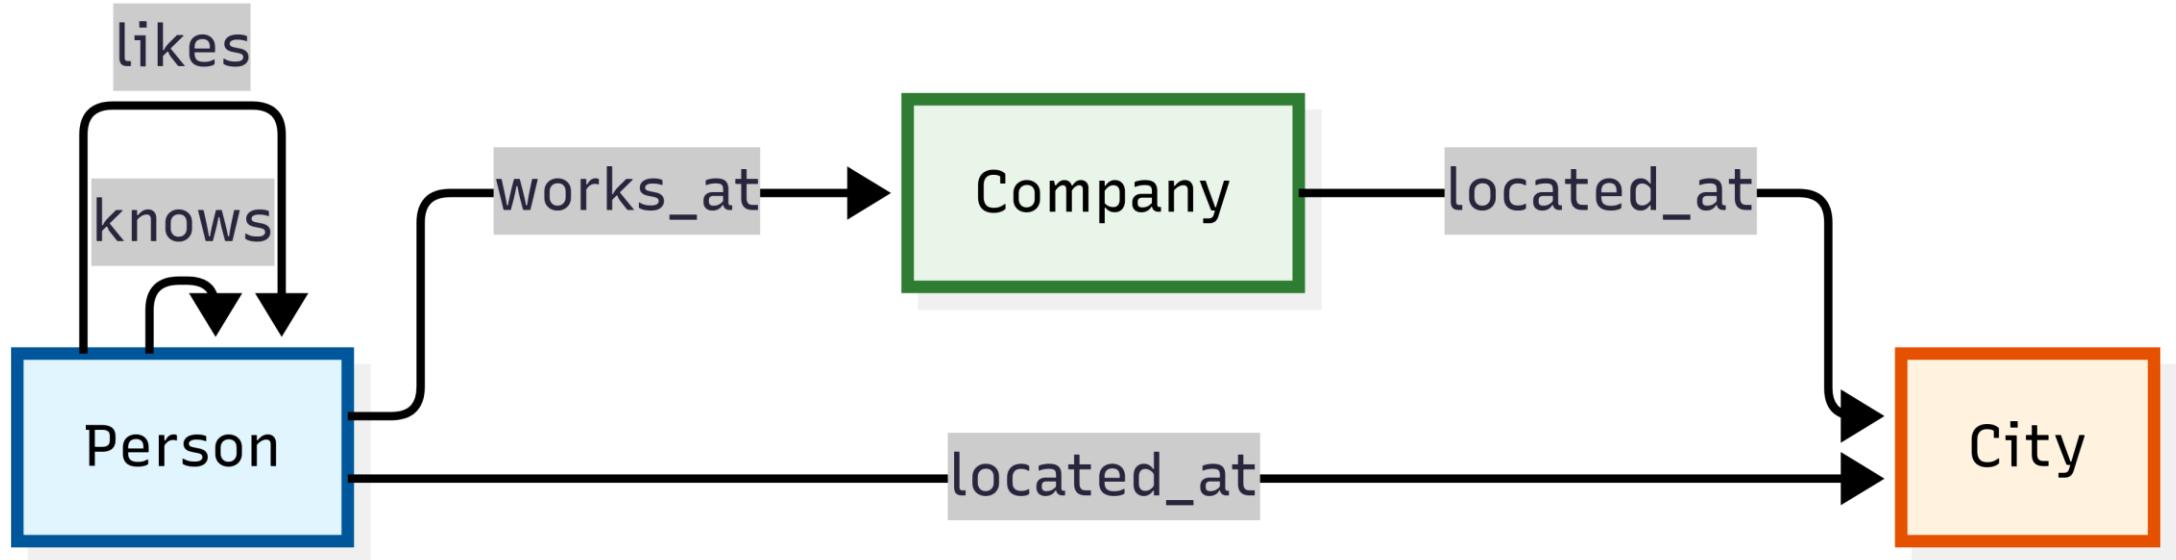
YellowHat 2026

-Henning Rauch-

Agenda

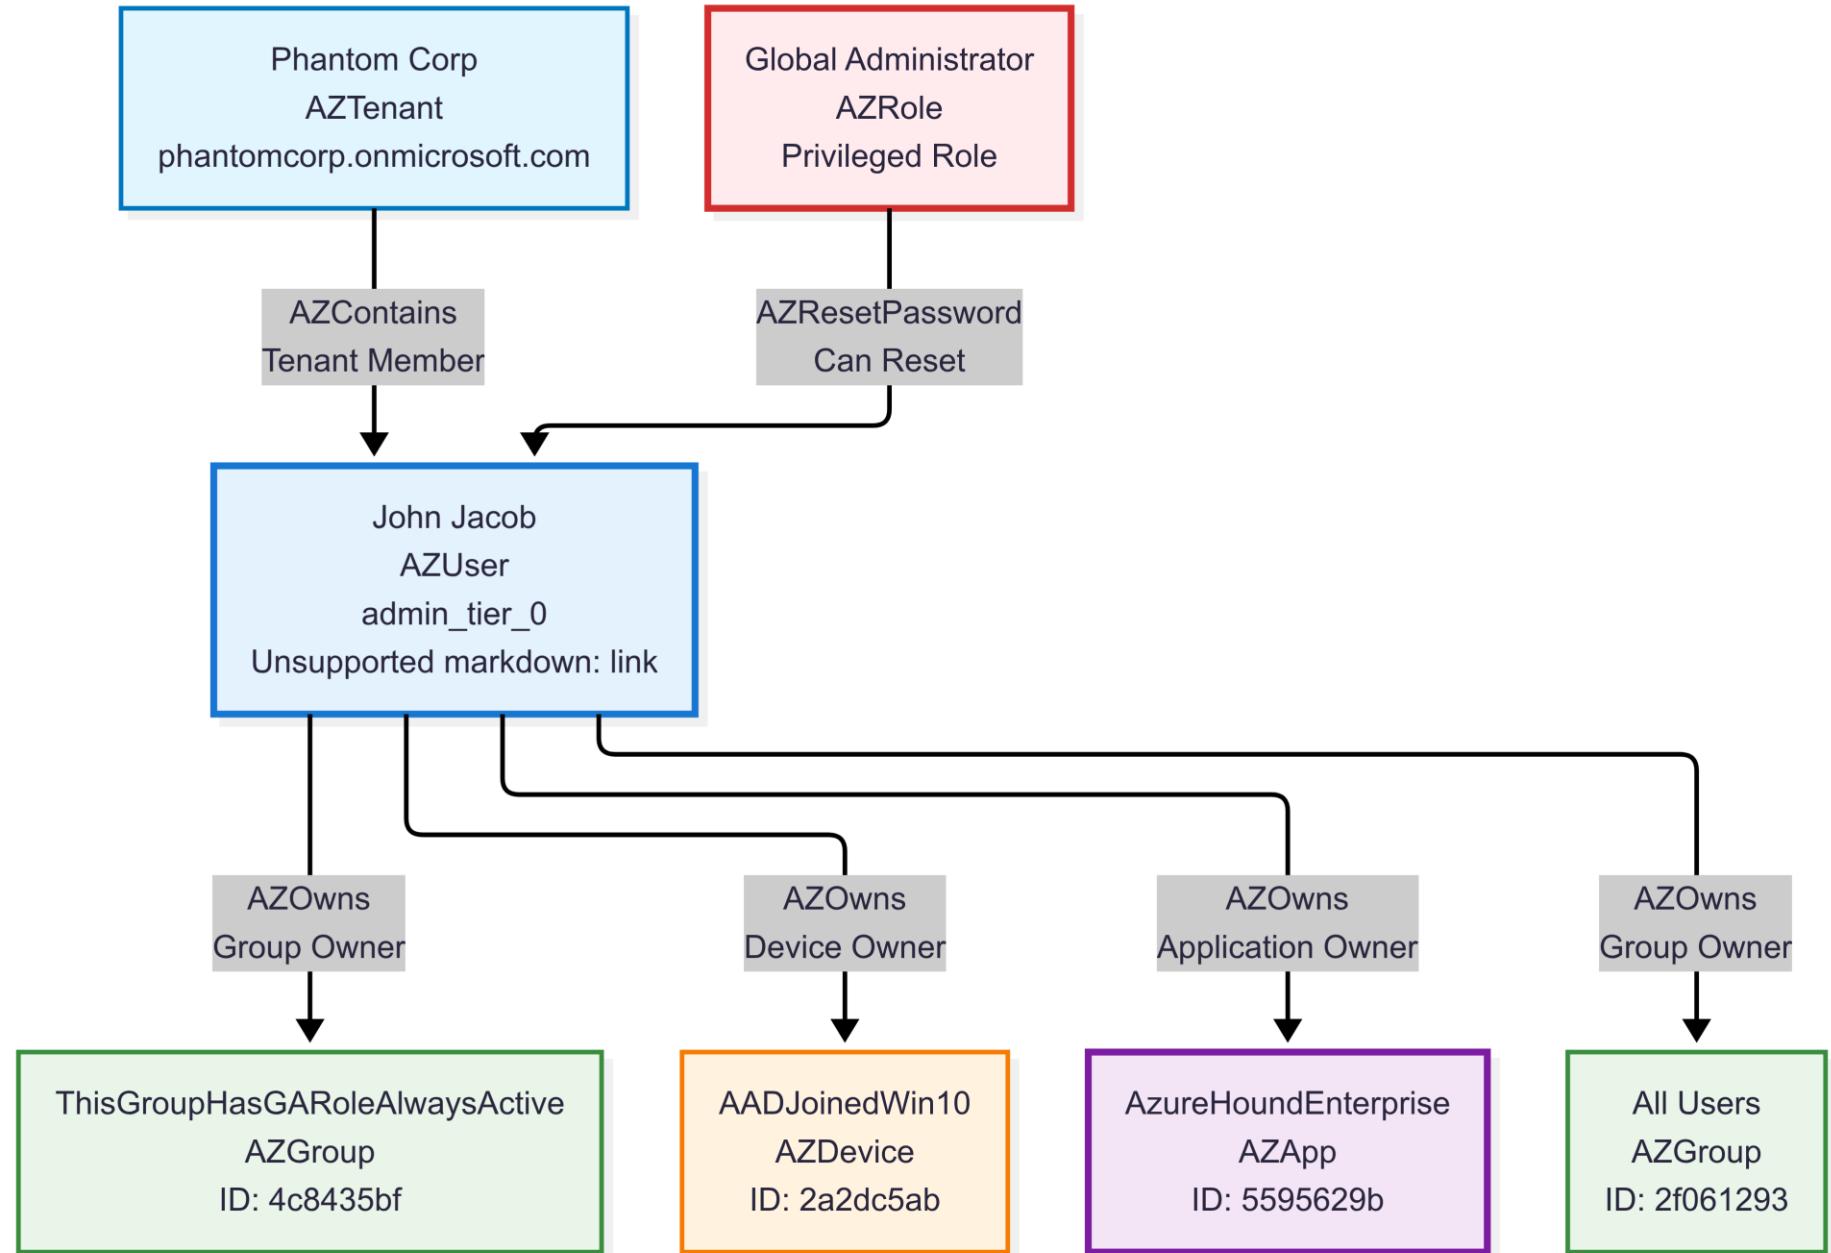
- Intro to graphs
- Graph services from Microsoft
- Deep dive into Kusto
- Demo

Graphs?



<https://learn.microsoft.com/en-us/kusto/query/graph-sample-data>

Graphs?



Kusto?

Size of Kusto

From Microsoft Digital Defense Report 2025

Our unique vantage point

Microsoft's global presence—spanning billions of users, millions of organizations, and a vast network of partners—provides us with an unparalleled perspective on the cybersecurity threat landscape.

100 trillion

100 trillion security signals processed daily

4.5 million

4.5 million net new malware file blocks every day

38 million

38 million identity risk detections analyzed in an average day

15,000+

15,000 partners in our security ecosystem, making it one of the largest in the world

34,000

34,000 full-time equivalent security engineers employed worldwide

5 billion

5 billion emails screened daily on average to protect users from malware and phishing

[Microsoft Digital Defense Report 2025 | Microsoft](#)

KustoCon
Learn | Share | Practice

Size of Kusto

From Fabric Realtime Intelligence slides...

502 T

Total messages per month

33 T

Streaming events processed monthly

12.5 EB

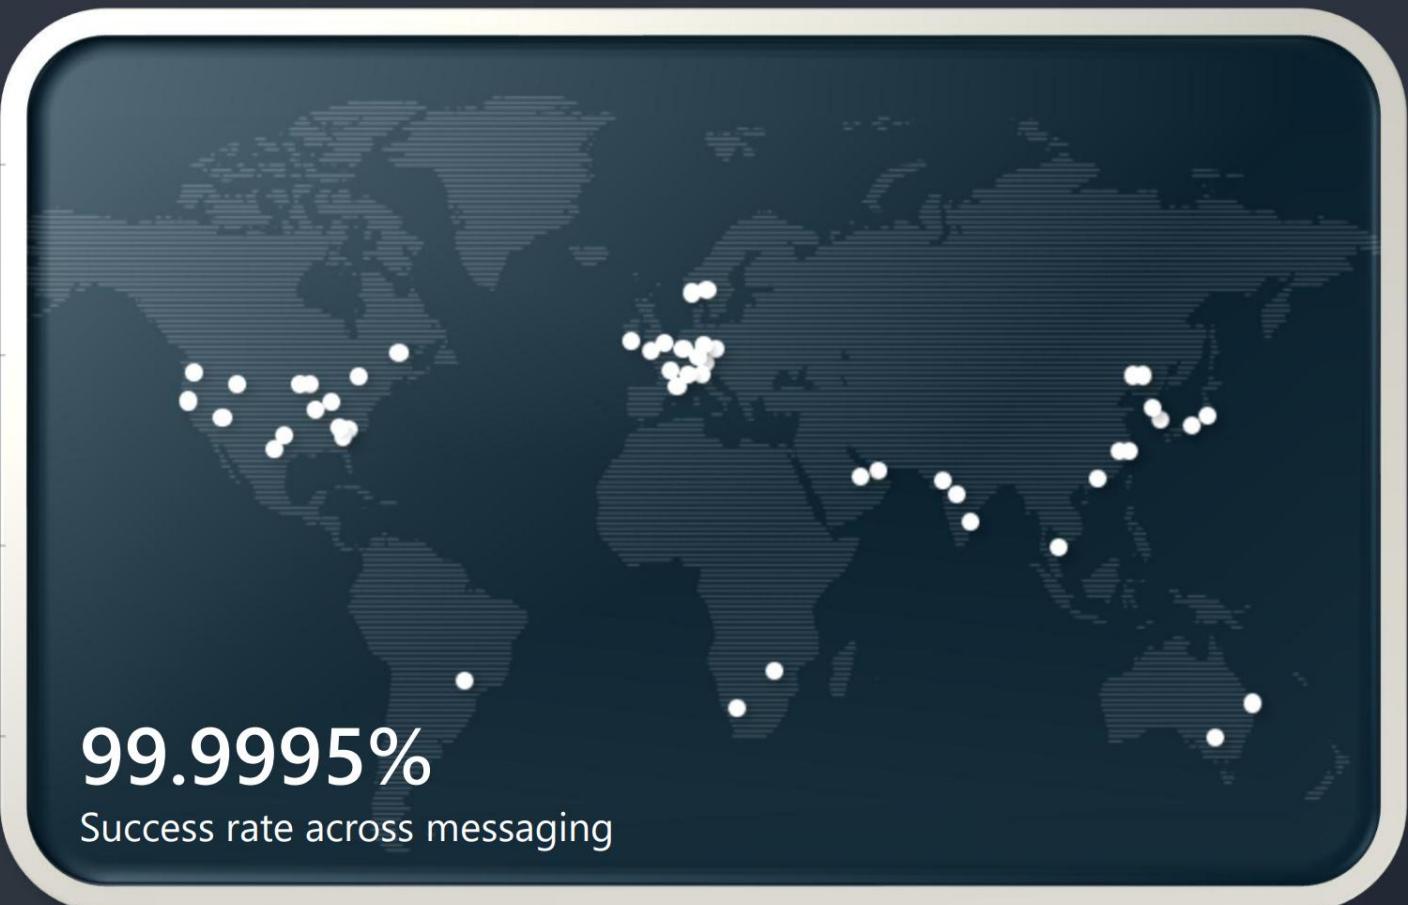
Events and logs per month

470 PB

Ingested daily

7 B

Real-time queries per day



KustoCon
Learn | Share | Practice

KustoCon

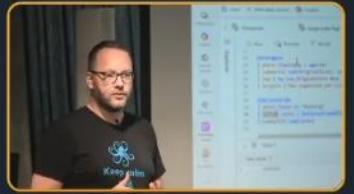
Learn | Share | Practice

The screenshot shows the Microsoft Fabric Data Explorer interface. On the left, there's a sidebar with various workspace options like Home, Workspaces, Copilot, OneLake Catalog, Monitor, Real-Time, and My workspace. The main area has tabs for Timeseries, Large scale logs, Microsoft scale, Engine features, Python, and a saved query named 'Features Demo'. Below these tabs is a toolbar with Run, Preview, Recall, Copy query, Save to Dashboard, KQL Tools, Export to CSV, Set alert, and More. The main content area displays a Kusto Query Language (KQL) query:

```
12 DataIngest
13 | where Timestamp > ago(5d)
14 | summarize sum(OriginalSize), sum(CompressedSize) by Source, bin(Timestamp, 1d)
15 | top 1 by sum_OriginalSize desc
16 | project ['Max ingestion per cluster on a single day'] = format_bytes(sum_OriginalSize), ['Compressed Size'] = format_bytes(sum_C
17
18 DimClustersMv
19 | where State == "Running"
20 | extend cores = GetCoresFromSKU(MachineSKU) * MachineCount
21 | summarize sum(cores)
22
```

Below the query, a table named 'Table 1' is shown with one row of data:

sum.cores
8,580,983



MORE VIDEOS

https://youtu.be/Lc-TgF9Oklg?si=VLBt5TbBrE1E3Z_i

Why Graphs for Cybersecurity?

- **The Problem with Tabular-Only Analysis**
 - Tables show events in isolation
 - Joins become complex and expensive at scale
 - Hard to visualize attack narratives
- **What Graphs Reveal**
 - Attack paths across multiple hops
 - Lateral movement chains
 - Blast radius of compromised accounts
 - Hidden dependencies between assets

The Microsoft Graph Technology Landscape

- CosmosDB (Gremlin)
- SQL databases
- Spark
- **Kusto**
 - **Azure Data Explorer**
 - **Eventhouse**
 - **(Log Analytics / Sentinel)**
- Fabric Graph (in preview)
- Microsoft Sentinel Graph



Fabric
Graph

Fabric RTI
Eventhouse

Azure Data
Explorer

Kusto

Microsoft's Graph Technologies: Choose Your Path

🚀 Kusto-based

(ADX, Eventhouse, Log
Analytics)

Scale-up graph

✓ Real-time streaming

✓ Breadth of analytics

🌐 Fabric Graph

Scale-out graph

GQL only

OneLake data (well aged)

🛡️ Sentinel Graph

KQL jobs (async, limited)

Security-focused UI

Incident, Hunting, Risk graphs

Details on Kusto



The cybersecurity world runs on KQL - and graph is just one capability among many.

Relational

Joins

Selection /
Projection

Aggregation

Timeseries

Functions

Anomaly
detection

Binning /
Windowing /
Forecasting

Geospatial

Grid based
spatial
indexing

Native
geospatial
functions

Sequence

Scan
operator

Predefined
patterns

Vector

Support for
high
dimensional
vectors

Similarity
functions

AI

Python
plugin

NL2KQL

MCP

Graph

Native graph
engine

Transient
and
persistent

GQL /
OpenCypher

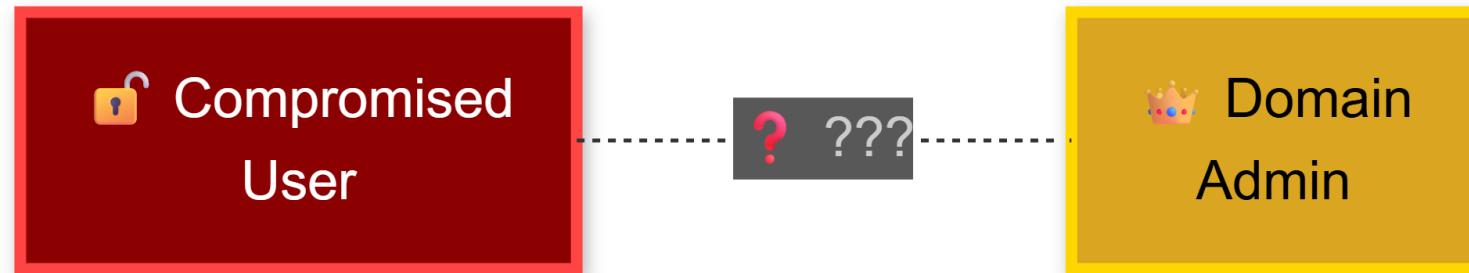
One language. One platform. All your security analytics.

Transient vs persistent graphs

Aspect	Transient	Persistent
Engine	Kusto	Kusto, Fabric Graph, Sentinel Graph
Storage	In-memory	Snapshot based storage
Lifespan	Query duration	Indefinite or until they are retained
Setup	None, Graph model optional	Graph model required
Best for	Exploration, ad-hoc	Repeated queries, dashboards, investigation on a fixed dataset
Scale limit	~10-100M elements (it depends)	Kusto: ~6B elements Fabric Graph: ? Sentinel Graph: ?
Refresh	Every query	snapshot refresh

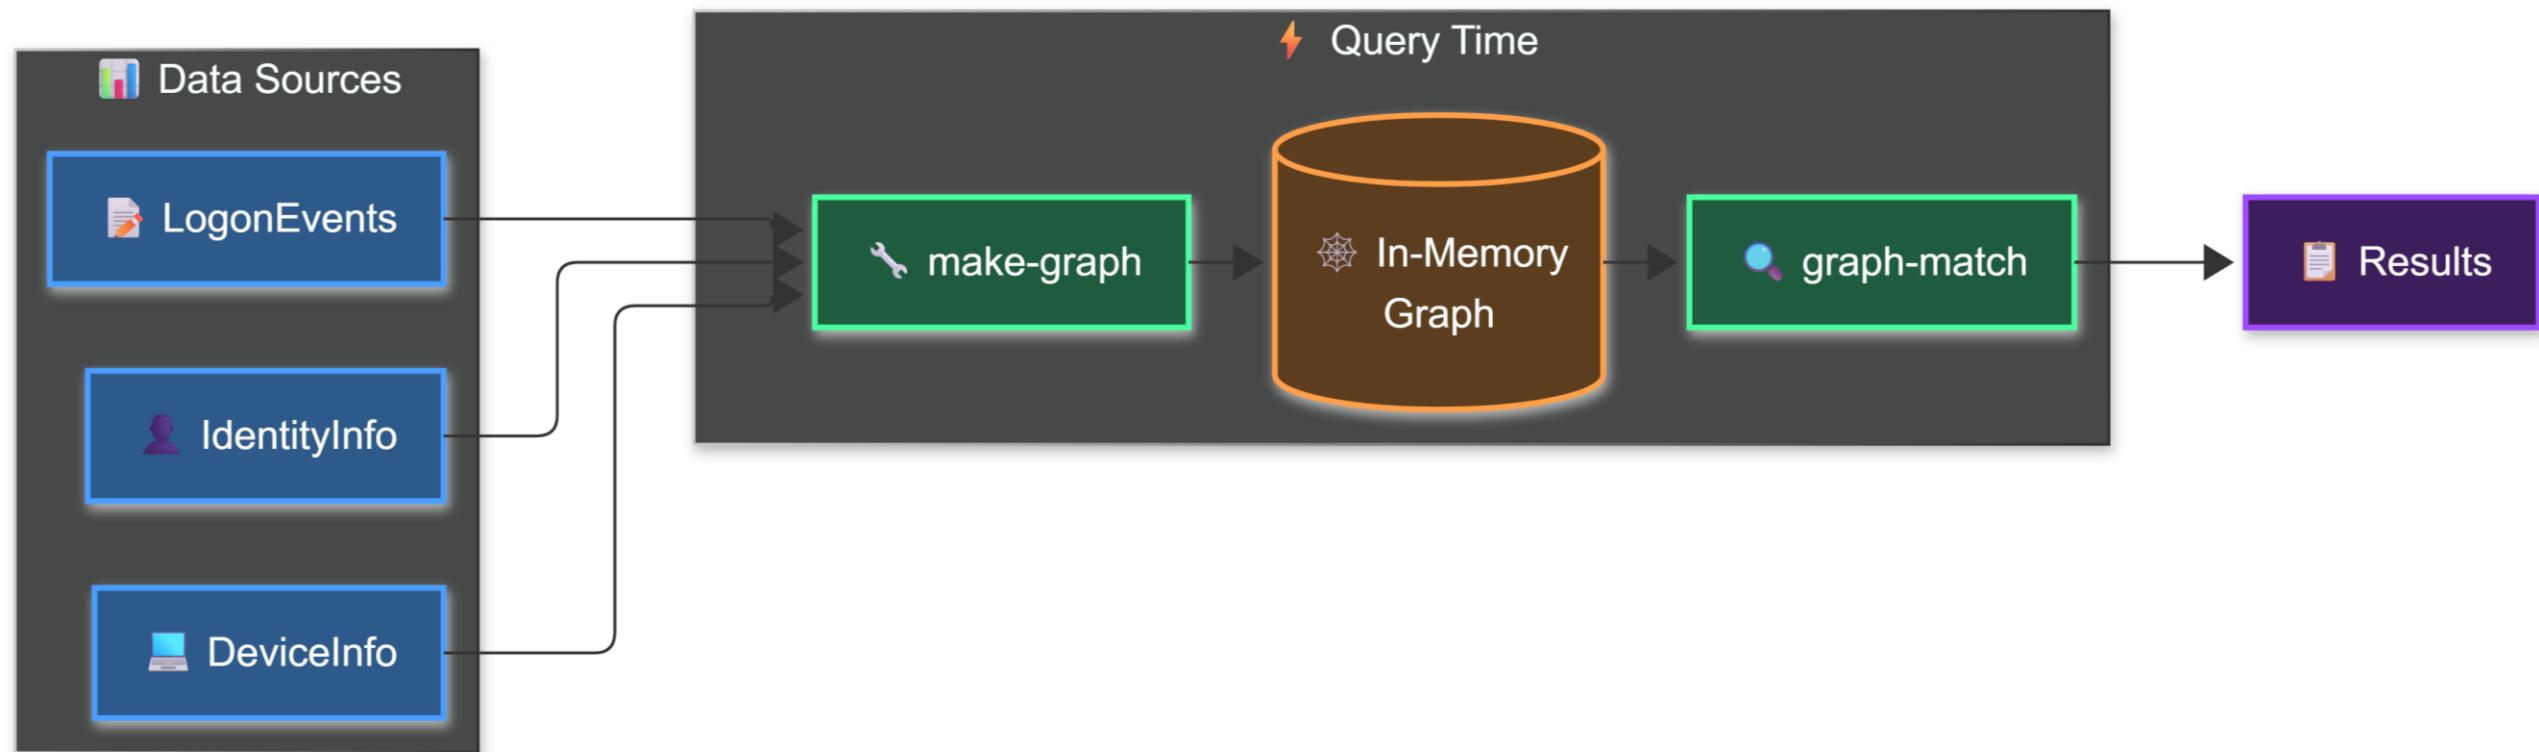
The Lateral Movement Scenario

- **Same problem, multiple solutions**
 - We want to detect attack paths from compromised users to Domain Controllers.
- **Data sources:**
 - IdentityInfo - User accounts and privileges
 - DeviceInfo - Machines and their roles
 - DeviceLogonEvents - Who logged in where, with what rights
- **Graph structure:**
 - Nodes: Identities and Devices
 - Edges: AdminTo, HasSession relationships



Approach 1: Transient Graph (make-graph)

- Graph built inline within the query
- No persistent storage
- Maximum flexibility
- Must rebuild for every execution

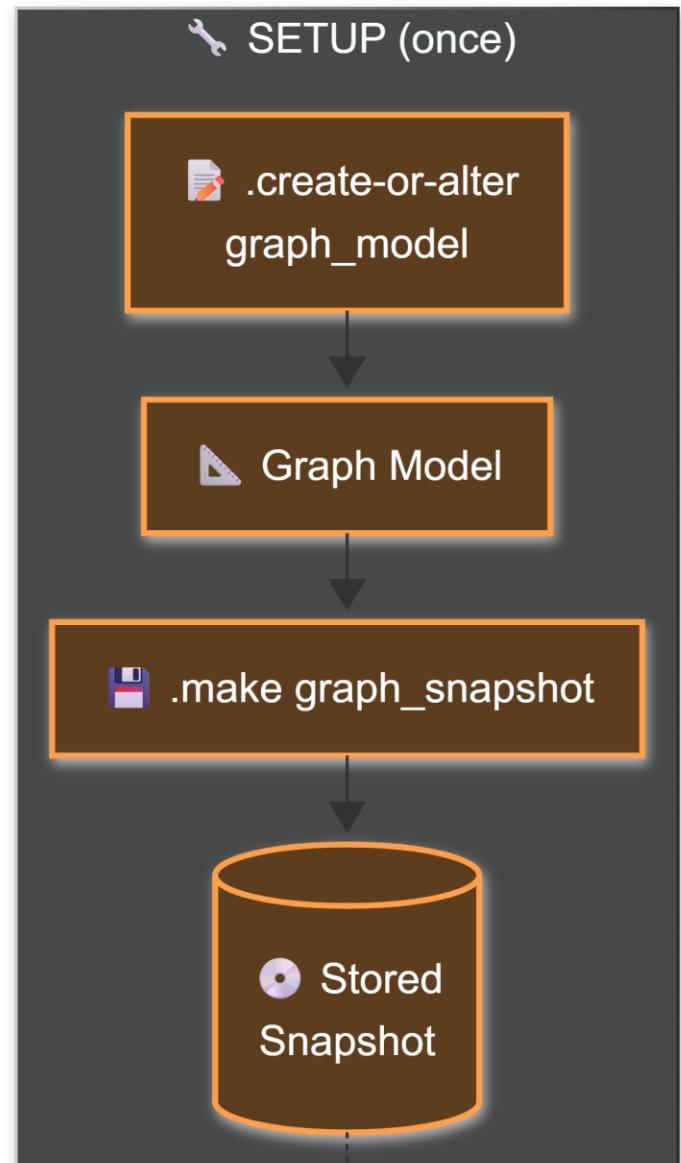


Approach 1: Transient Graph (make-graph)

```
// Transient lateral movement detection
let identities =
    IdentityInfo
    | summarize arg_max(Timestamp, *) by AccountSid
    | project
        NodeId = AccountSid,
        NodeLabel = "Identity",
        AccountName,
        IsAdmin = IsDomainAdmin or IsLocalAdmin;
let devices =
    DeviceInfo
    | summarize arg_max(Timestamp, *) by DeviceId
    | project
        NodeId = DeviceId,
        NodeLabel = "Device",
        DeviceName,
        IsDomainController;
let nodes = union identities, devices;
```

Approach 2: Persistent Graph Model

- Schema and structure defined once
- Snapshots materialize the graph
- Query against pre-built snapshots
- Separates definition from execution



Approach 2: Create the Graph Model

```
.create-or-alter graph_model LateralMovement """
{
    "Schema": {
        "Nodes": {
            "Identity": {
                "AccountSid": "string",
                "AccountName": "string",
                "IsAdmin": "bool"
            },
            "Device": {
                "DeviceId": "string",
                "DeviceName": "string",
                "IsDomainController": "bool"
            }
        }
    }
}
```

Approach 2: (continued) Definition Steps

```
"Definition": {  
    "Steps": [  
        {  
            "Kind": "AddNodes",  
            "Query": "IdentityInfo | summarize arg_max(Timestamp, *) by AccountSid | project NodeId=AccountSid, AccountName, IsAdmin=IsDomainAdmin or IsLocalAdmin",  
            "NodeIdColumn": "NodeId",  
            "Labels": ["Identity"]  
        },  
        {  
            "Kind": "AddNodes",  
            "Query": "DeviceInfo | summarize arg_max(Timestamp, *) by DeviceId | project NodeId=DeviceId, DeviceName, IsDomainController",  
            "NodeIdColumn": "NodeId",  
            "Labels": ["Device"]  
        },  
        {  
            "Kind": "AddEdges",  
            "Query": "DeviceLogonEvents | where IsLocalAdmin == true | project SourceNode=AccountId, TargetNode=DeviceId, Timestamp",  
            "SourceColumn": "SourceNode",  
            "TargetColumn": "TargetNode",  
            "Labels": ["AdminTo"]  
        },  
        {  
            "Kind": "AddEdges",  
            "Query": "DeviceLogonEvents | where IsLocalAdmin == false | project SourceNode=AccountId, TargetNode=DeviceId, Timestamp",  
            "SourceColumn": "SourceNode",  
            "TargetColumn": "TargetNode",  
            "Labels": ["HasSession"]  
        }  
    ]  
}
```

Approach 2: Create a Graph Snapshot

```
// Create a snapshot (synchronous - waits for completion)
.make graph_snapshot DailySnapshot_20260109 from LateralMovement
```

Returns:

Name	SnapshotTime	NodesCount	EdgesCount
DailySnapshot_20260109	2026-01-09T06:00:00Z	45,823	1,247,891

```
// For large graphs, use async
.make async graph_snapshot WeeklySnapshot from LateralMovement
```

Approach 2: Query the Persistent Graph

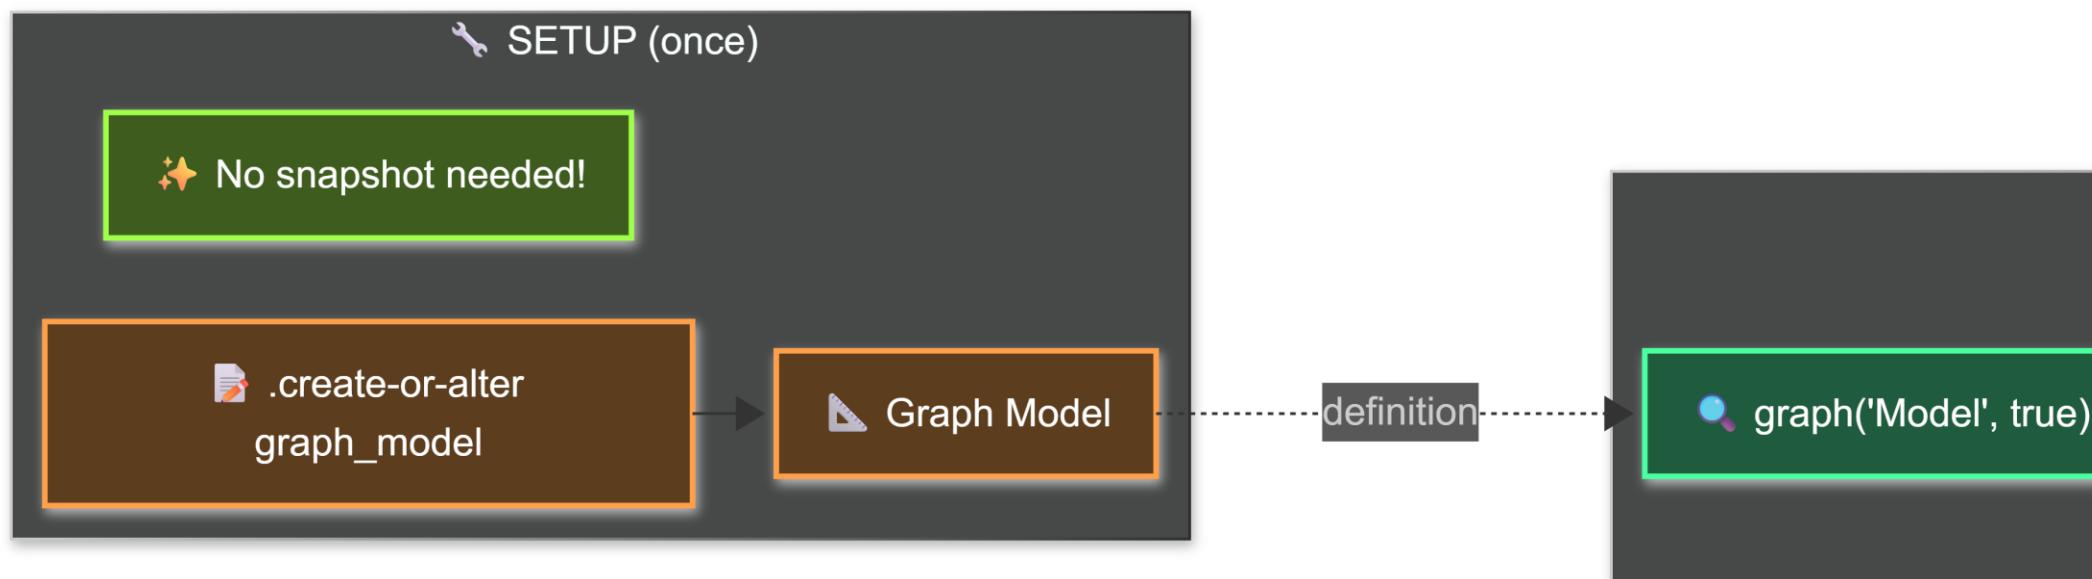
```
// Query the latest snapshot
graph("LateralMovement")
| graph-match (user)-[path*1..5]->(target)
  where labels(user) has "Identity"
    and labels(target) has "Device"
    and user.IsAdmin == false
    and target.IsDomainController == true
  project
    StartUser = user.AccountName,
    PathLength = array_length(path),
    TargetDC = target.DeviceName
  | order by PathLength asc
```

Approach 2: Query the Persistent Graph

```
// Query a specific snapshot (for historical comparison)
graph("LateralMovement", "DailySnapshot_20260108")
| graph-match (user)-[path*1..5]->(target)
  where labels(user) has "Identity"
    and labels(target) has "Device"
    and user.IsAdmin == false
    and target.IsDomainController == true
  project StartUser = user.AccountName, PathLength = array_length(path)
```

Approach 3: Transient Graph from Model

- Uses the model's schema and definition
- Builds the graph fresh at query time (like make-graph)
- No snapshot required
- Always uses current data



Approach 3: Transient Graph from Model

```
// Transient graph from model - always uses latest data
graph("LateralMovement", true)
| graph-match (user)-[path*1..5]->(target)
  where labels(user) has "Identity"
    and labels(target) has "Device"
    and user.IsAdmin == false
    and target.IsDomainController == true
project
  StartUser = user.AccountName,
  PathLength = array_length(path),
  TargetDC = target.DeviceName
| order by PathLength asc
```

Comparing the query syntax

Approach	Code
Transient from make-graph	edges make-graph Src --> Tgt with nodes on Id
Snapshot	graph("Model") or graph("Model", "SnapshotName")
Transient from model	graph("Model", true)

You can run all graph operators (graph-match, graph-to-table) on all type of Kusto graphs regardless of how they have been created.

Three approaches compared

Aspect	make-graph	graph(Model, true)	graph(Model)
Setup	None	Create model	Model + snapshots
Schema	Inline	Defined once	Defined once
Data freshness	Real-time	Real-time	Snapshot time
Query performance	Rebuild each time	Rebuild each time	Pre-built
Historical	No	No	Yes
Maintenance	None	Model-only	Model + snapshots
Best for	Ad-hoc exploration	Structured real-time	Production dashboards, Fix investigations

Three approaches compared

🔧 FLEXIBILITY

⚡ make-graph

No structure
Real-time

⚖️ BALANCED

⟳ graph(M,true)

Structured
Real-time

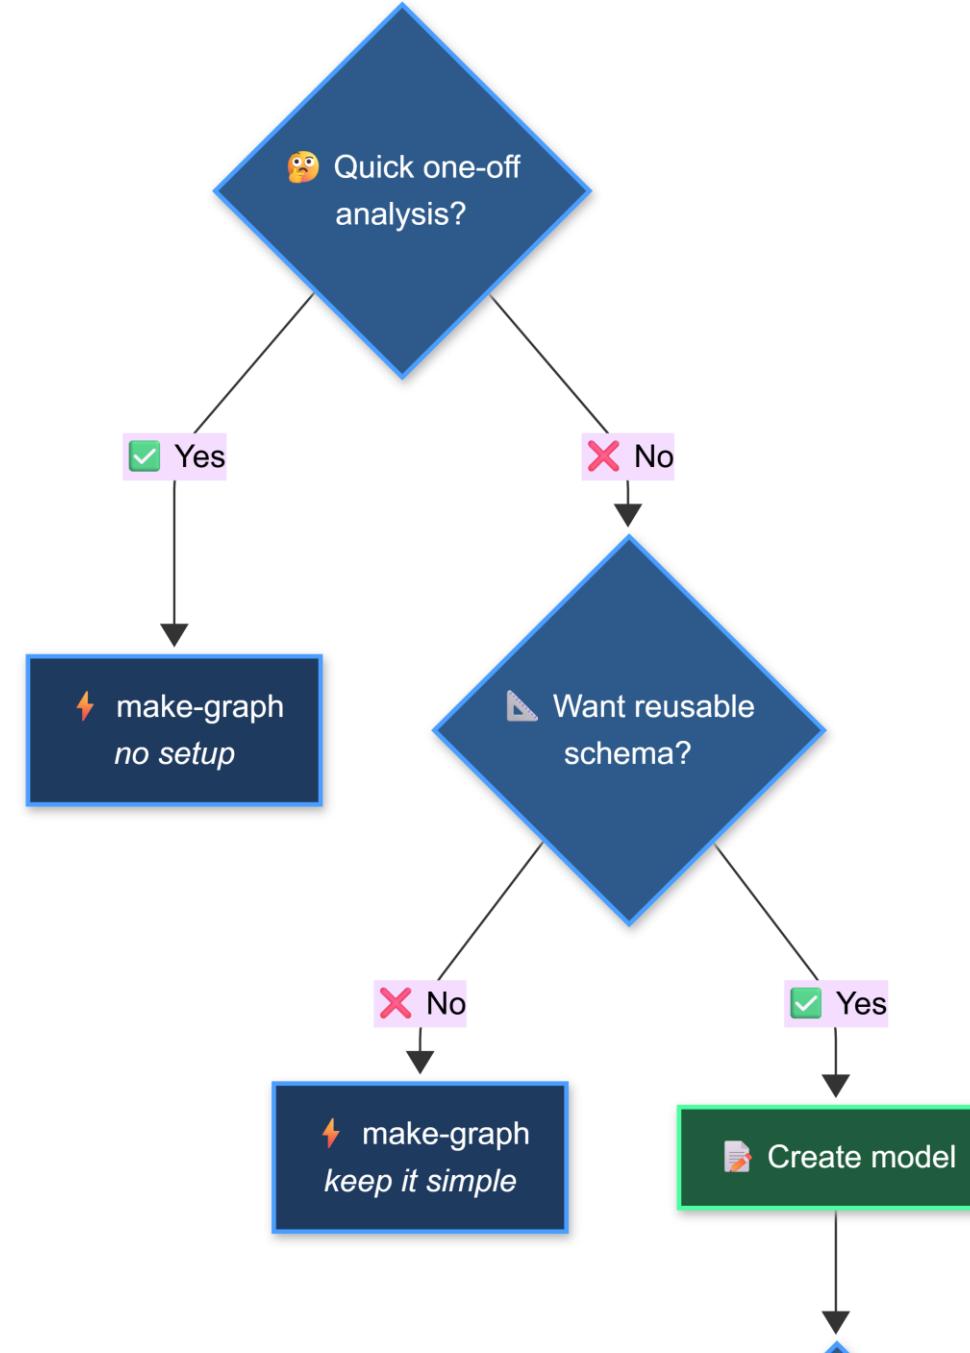
🚀 PERFORMANCE

💾 graph(M)

Structured
Point-in-time



Choosing Your Approach: Decision Tree



Cybersecurity Functions

Cybersecurity Functions: Multiple Algebras Required

- Graph functions

Function	Purpose	Question answered
graph_blast_radius_fl()	Downstream impact	"If this is compromised, what's at risk?"
graph_exposure_perimeter_fl()	Upstream exposure	"What could compromise this asset?"
graph_node_centrality_fl()	Critical nodes	"Which nodes are chokepoints?"
graph_path_discovery_fl()	Path enumeration	"How can attacker reach target?"

- Anomaly detection function

Function	Technique	Use Case
detect_anomalous_access_cf_fl()	Collaborative filtering	Unusual user-resource access
detect_anomalous_new_entity_fl()	Poisson distribution	New suspicious entities
detect_anomalous_spike_fl()	Z-score / Q-score	Data exfiltration spikes

Demo time

- 1) Find attack paths using exposure graph tables
- 2) Combining Algebras: Spike + Blast Radius
- 3) Collaborative filtering anomaly detection + Graph exposure perimeter = Prioritized resource protection
- 4) Graph-RAG for Threat Intelligence

Key takeaways

- **Multiple Microsoft graph technologies exist**
 - Kusto, Fabric Graph, Sentinel Graph - each with pros and cons
 - Choose based on your scale, real-time needs, and workflow
- **One algebra is not enough for cybersecurity**
 - Graph alone won't protect you
 - Combine graph with statistics, ML, time series analysis
- **KQL: One language, all algebras, one query**
- **Use what Microsoft already built**
 - Cybersecurity functions ready to use

Next Steps

- [MCP Registry | Fabric Real-Time Intelligence](#)
 - Go check out MCP server for KQL
- [langgraph-kusto · PyPI](#)
 - Kusto-backed storage and checkpointing for LangGraph
 - Enables multi-agent systems with human in the loop

Thank you!

YELLOWHAT

HARD HATS FOR CYBER THREATS

Elite

VirtualMetric

Gold



glueck■kanja



bnevinn

baseVISION

water
IT Security & Defense



Rubicon

Bronze



SITS

TRAXION

MOSADEX E-HEALTH



valid

INSPARK

yellow arrow

DELTA-N

ENDOR LABS

yellowhat.live

